

Monitoring and Instrumentation

There are several ways to monitor Spark applications: web UIs, metrics, and external instrumentation.

Web Interfaces

Every `SparkContext` launches a web UI, by default on port 4040, that displays useful information about the application. This includes:

- A list of scheduler stages and tasks
- A summary of RDD sizes and memory usage
- Environmental information.
- Information about the running executors

You can access this interface by simply opening `http://<driver-node>:4040` in a web browser. If multiple `SparkContexts` are running on the same host, they will bind to successive ports beginning with 4040 (4041, 4042, etc).

Note that this information is only available for the duration of the application by default. To view the web UI after the fact, set `spark.eventLog.enabled` to `true` before starting the application. This configures Spark to log Spark events that encode the information displayed in the UI to persisted storage.

Viewing After the Fact

Spark's Standalone Mode cluster manager also has its own [web UI](#). If an application has logged events over the course of its lifetime, then the Standalone master's web UI will automatically re-render the application's UI after the application has finished.

If Spark is run on Mesos or YARN, it is still possible to reconstruct the UI of a finished application through Spark's history server, provided that the application's event logs exist. You can start the history server by executing:

```
./sbin/start-history-server.sh
```

When using the file-system provider class (see `spark.history.provider` below), the base logging directory must be supplied in the `spark.history.fs.logDirectory` configuration option, and should contain sub-directories that each represents an application's event logs. This creates a web interface at `http://<server-url>:18080` by default. The history server can be configured as follows:

Environment Variable	Meaning
SPARK_DAEMON_MEMORY	Memory to allocate to the history server (default: 512m).
104.130.219.184	JVM options for the history server (default: none).
SPARK_PUBLIC_DNS	The public address for the history server. If this is not set, links to application history may use the internal address of the server, resulting in broken links (default: none).
SPARK_HISTORY_OPTS	<code>spark.history.*</code> configuration options for the history server (default: none).

Property Name	Default	Meaning
spark.history.provider	org.apache.spark.deploy.history.FsHistoryProvider	Name of the class implementing the application history backend. Currently there is only one implementation, provided by Spark,

which looks for application logs stored in the file system.

<code>spark.history.fs.logDirectory</code>	<code>file:/tmp/spark-events</code>	Directory that contains application event logs to be loaded by the history server
<code>spark.history.fs.updateInterval</code>	10	The period, in seconds, at which information displayed by this history server is updated. Each update checks for any changes made to the event logs in persisted storage.
<code>spark.history.retainedApplications</code>	50	The number of application UIs to retain. If this cap is exceeded, then the oldest applications will be removed.
<code>spark.history.ui.port</code>	18080	The port to which the web interface of the history server binds.
<code>spark.history.kerberos.enabled</code>	false	Indicates whether the history server should use kerberos to login. This is useful if the history server is accessing HDFS files on a secure Hadoop cluster. If this is true, it uses the configs <code>spark.history.kerberos.principal</code> and <code>spark.history.kerberos.keytab</code> .
<code>spark.history.kerberos.principal</code>	(none)	Kerberos principal name for the History Server.
<code>spark.history.kerberos.keytab</code>	(none)	Location of the kerberos keytab file for the History Server.
<code>spark.history.ui.acls.enable</code>	false	Specifies whether acls should be checked to authorize users viewing the applications. If enabled, access control checks are made regardless of what the individual application had set for <code>spark.ui.acls.enable</code> when the application was run. The application owner will always have authorization to view their own application and any users specified via <code>spark.ui.view.acls</code> when the application was run will also have authorization to view that application. If disabled, no access control checks are made.

Note that in all of these UIs, the tables are sortable by clicking their headers, making it easy to identify slow tasks, data skew, etc.

Note that the history server only displays completed Spark jobs. One way to signal the completion of a Spark job is to stop

the Spark Context explicitly (`sc.stop()`), or in Python using the `with SparkContext()` as `sc`: to handle the Spark Context setup and tear down, and still show the job history on the UI.

Metrics

Spark has a configurable metrics system based on the [Coda Hale Metrics Library](#). This allows users to report Spark metrics to a variety of sinks including HTTP, JMX, and CSV files. The metrics system is configured via a configuration file that Spark expects to be present at `$SPARK_HOME/conf/metrics.properties`. A custom file location can be specified via the `spark.metrics.conf` [configuration property](#). Spark's metrics are decoupled into different *instances* corresponding to Spark components. Within each instance, you can configure a set of sinks to which metrics are reported. The following instances are currently supported:

- `master`: The Spark standalone master process.
- `applications`: A component within the master which reports on various applications.
- `worker`: A Spark standalone worker process.
- `executor`: A Spark executor.
- `driver`: The Spark driver process (the process in which your `SparkContext` is created).

Each instance can report to zero or more *sinks*. Sinks are contained in the `org.apache.spark.metrics.sink` package:

- `ConsoleSink`: Logs metrics information to the console.
- `CSVSink`: Exports metrics data to CSV files at regular intervals.
- `JmxSink`: Registers metrics for viewing in a JMX console.
- `MetricsServlet`: Adds a servlet within the existing Spark UI to serve metrics data as JSON data.
- `GraphiteSink`: Sends metrics to a Graphite node.

Spark also supports a Ganglia sink which is not included in the default build due to licensing restrictions:

- `GangliaSink`: Sends metrics to a Ganglia node or multicast group.

To install the `GangliaSink` you'll need to perform a custom build of Spark. ***Note that by embedding this library you will include LGPL-licensed code in your Spark package.*** For sbt users, set the `SPARK_GANGLIA_LGPL` environment variable before building. For Maven users, enable the `-Pspark-ganglia-lgpl` profile. In addition to modifying the cluster's Spark build user applications will need to link to the `spark-ganglia-lgpl` artifact.

The syntax of the metrics configuration file is defined in an example configuration file, `$SPARK_HOME/conf/metrics.properties.template`.

Advanced Instrumentation

Several external tools can be used to help profile the performance of Spark jobs:

- Cluster-wide monitoring tools, such as [Ganglia](#), can provide insight into overall cluster utilization and resource bottlenecks. For instance, a Ganglia dashboard can quickly reveal whether a particular workload is disk bound, network bound, or CPU bound.
- OS profiling tools such as [dstat](#), [iostat](#), and [iotop](#) can provide fine-grained profiling on individual nodes.
- JVM utilities such as `jstack` for providing stack traces, `jmap` for creating heap-dumps, `jstat` for reporting time-series statistics and `jconsole` for visually exploring various JVM properties are useful for those comfortable with JVM internals.